



# DEEP LEARNING

---

# AGENDA

---

- INTRODUCTION TO DEEP LEARNING
- WHY DEEP LEARNING?
- WHY DEEP LEARNING NOW?
- INTRODUCTION TO CNN
- BASICS OF CNN
- DIGIT RECOGNITION USING LENET





# **ARTIFICIAL INTELLIGENCE**

Programs with the ability to learn and reason like humans

## **MACHINE LEARNING**

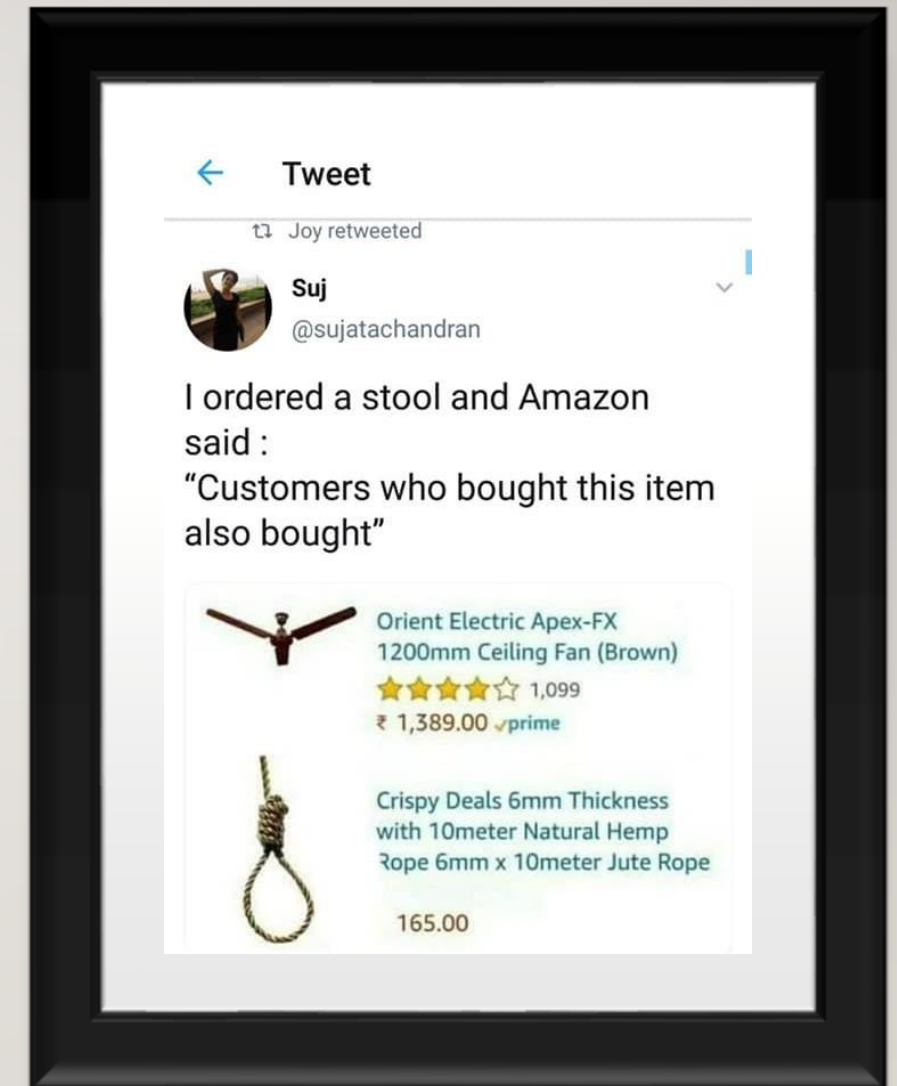
Algorithms with the ability to learn without being explicitly programmed

## **DEEP LEARNING**

Subset of machine learning in which artificial neural networks adapt and learn from vast amounts of data

# MACHINE LEARNING USE CASE

- Amazon uses machine learning to improve user experiences and recommendation based on their usage
- Netflix uses machine learning to suggest content to users




# WHY DEEP LEARNING?

---



Speech  
Recognition




Computer  
Vision



Natural Language  
Processing

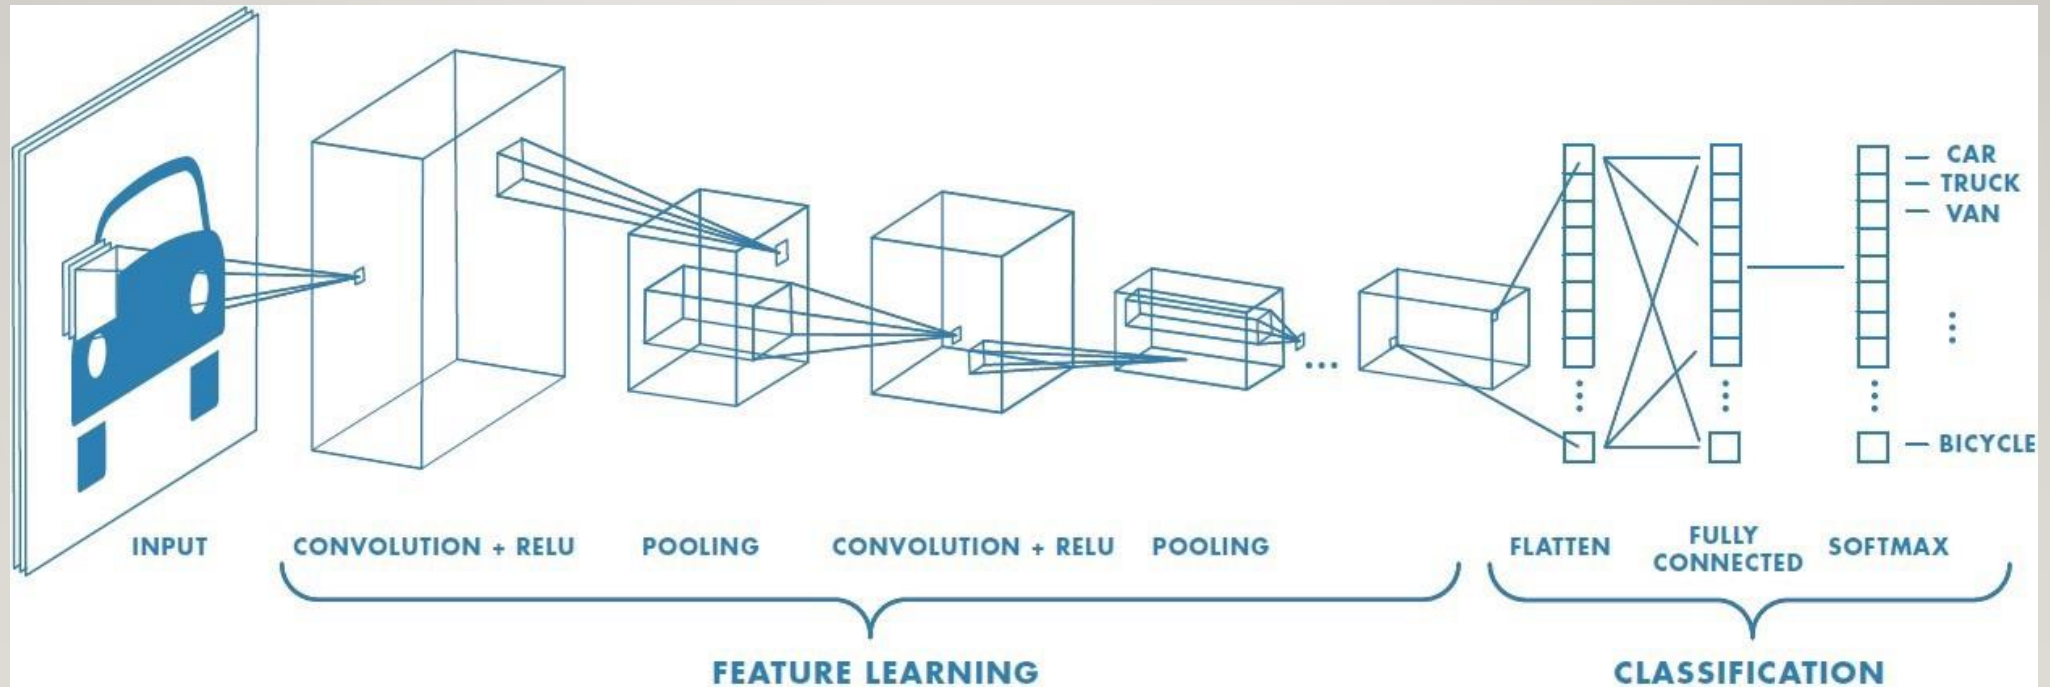
# WHY DEEP LEARNING NOW?


---





# TYPICAL CONVOLUTIONAL NEURAL NETWORK





A more accepted convention,

$$y = 1 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i < 0$$

where,  $x_0 = 1$  and  $w_0 = -\theta$

$x_1$	$x_2$	OR	
0	0	0	$w_0 + \sum_{i=1}^2 w_i x_i < 0$
1	0	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$
0	1	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$
1	1	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$

$$w_0 + w_1 \cdot 0 + w_2 \cdot 0 < 0 \Rightarrow w_0 < 0$$


$$w_0 + w_1 \cdot 0 + w_2 \cdot 1 \geq 0 \Rightarrow w_2 > -w_0$$

$$w_0 + w_1 \cdot 1 + w_2 \cdot 0 \geq 0 \Rightarrow w_1 > -w_0$$

$$w_0 + w_1 \cdot 1 + w_2 \cdot 1 \geq 0 \Rightarrow w_1 + w_2 > -w_0$$

One possible solution is

$$w_0 = -1, w_1 = 1.1, w_2 = 1.1$$



## INTRODUCTION TO CNN: PERCEPTRON




## PERCEPTRON LEARNING ALGORITHM

---

```
P ← inputs with label 1;  
N ← inputs with label 0;  
Initialize  $\mathbf{w} = [w_1, w_2, \dots, w_N]$  randomly;  
while !convergence do  
    Pick random  $\mathbf{x} \in P \cup N$  ;  
    if  $\mathbf{x} \in P$  and  $\sum_{i=0}^n w_i * x_i < 0$  then  
        |  $\mathbf{w} = \mathbf{w} + \mathbf{x}$  ;  
    end  
    if  $\mathbf{x} \in N$  and  $\sum_{i=0}^n w_i * x_i \geq 0$  then  
        |  $\mathbf{w} = \mathbf{w} - \mathbf{x}$  ;  
    end  
end  
//the algorithm converges when all the  
inputs are classified correctly
```

Input		Output
x	y	$x \oplus y$
0	0	0
0	1	1
1	0	1
1	1	0



# LIMITATIONS OF PERCEPTRON

A close-up shot from the movie Inception showing Leonardo DiCaprio and Inez. Leonardo DiCaprio is on the left, looking slightly to the right with a serious expression. Inez is on the right, seen in profile, looking towards Leonardo. The background is blurred, showing what appears to be a modern interior with large windows.

**WE NEED TO GO**

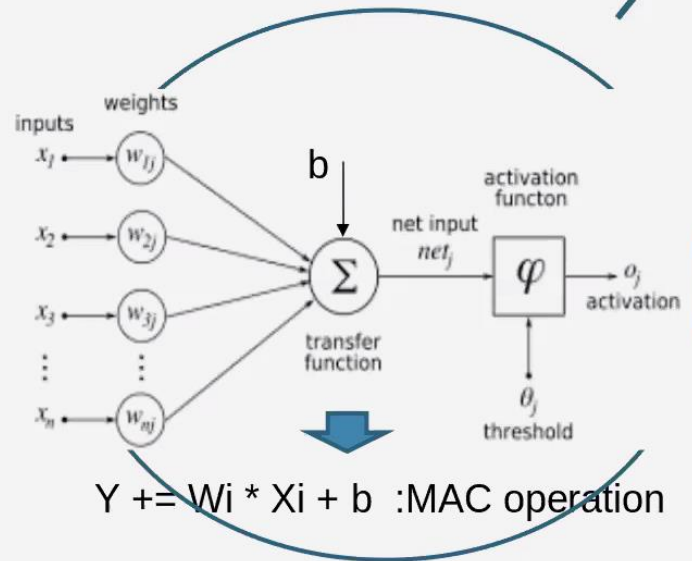
**DEEPER**



# CNN - Neural Net classifier




Single Neuron architecture

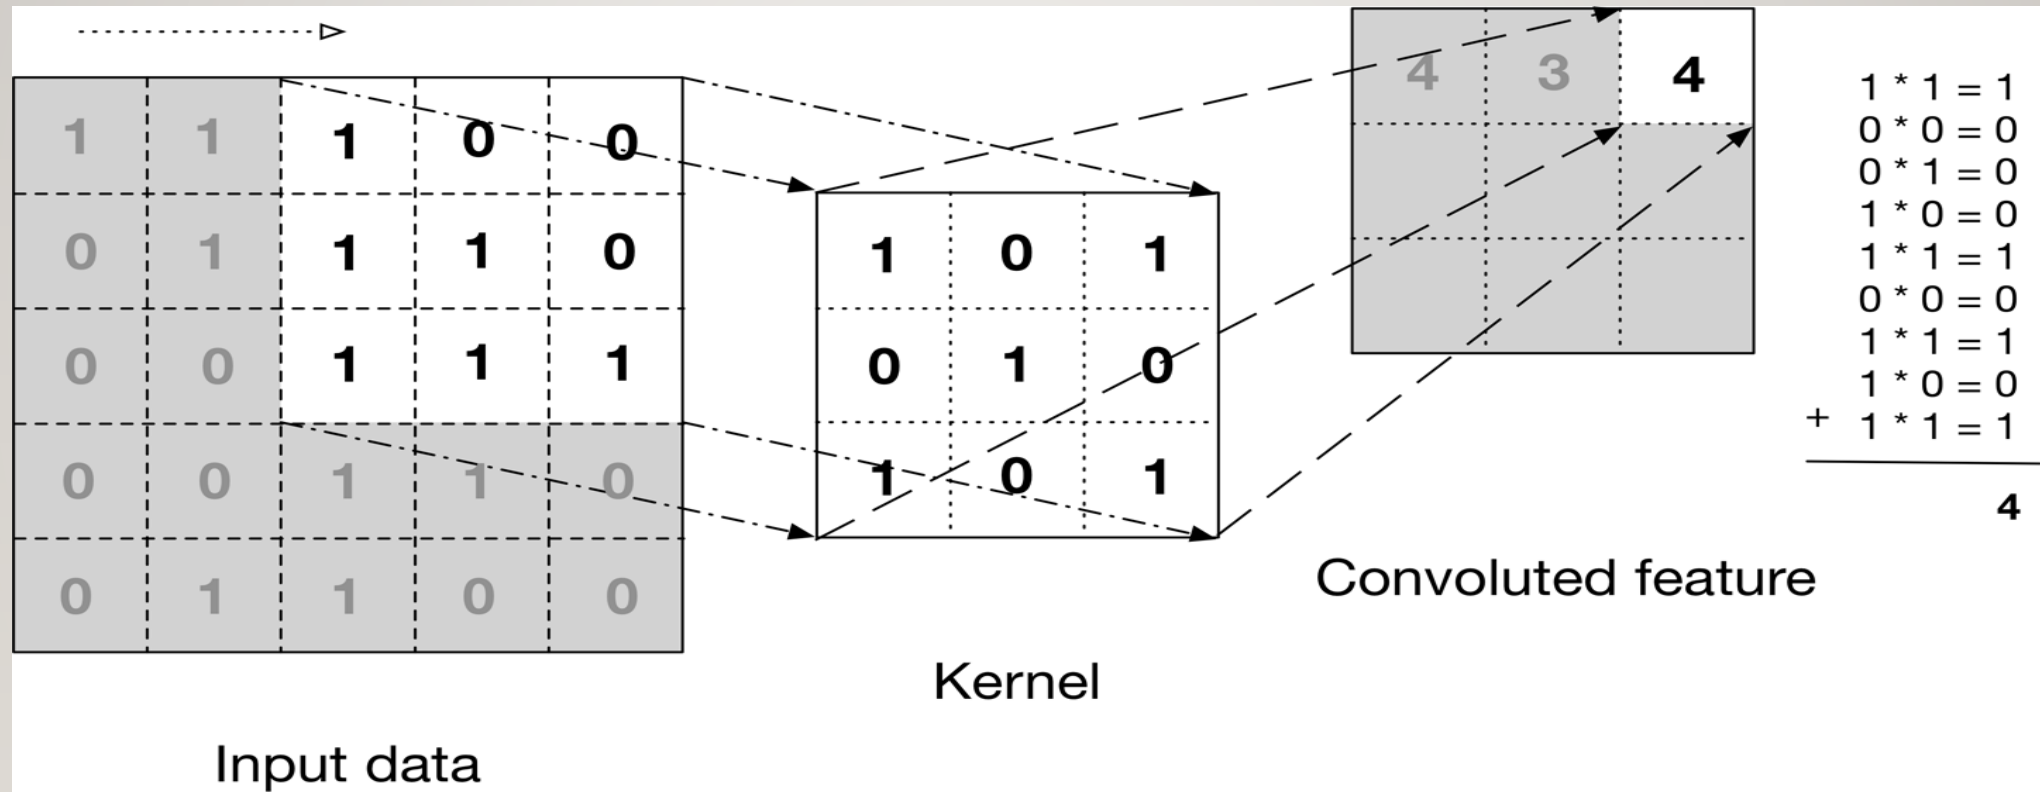


Activation function = Nonlinear function  $\Rightarrow$  exp, tanh etc

MLP: Fully Connected Neural Net



# CONVOLUTION



# TERMINOLOGY IN CNN

---

1. Kernel/filter/Weights: Learnable parameters of our model
2. Stride: how many cells to move
3. Bias: how far off our predictions are from real values
4. Feature Maps: output of Intermediator layers




# Vertical edge detection

↓ ↓

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0

6x6



\*

1	0	-1
1	0	-1
1	0	-1


3x3

=


↓

0	30	30	0
0	30	30	0
0	30	30	0
0	<u>30</u>	30	0


↑ 4x4



\*



# POOLING: DOWN SAMPLING




Scan Right by H-stride

1. Select max out of 4 pixel in Max Pooling
2. Select average of all 4 pixel in Average Pooling
3. Size is typically (2 x 2) but it can be any integer
4. In (2 x 2) stride is 2
5. Pooling reduces computations in CNN model and handles dimensionality problem

Scan Down by V-stride

# NEAREST UP SAMPLING POOLING

## 2x2 Nearest Up-sampling



Scan Down by V-stride




Scan Right by H-stride

1. Replicate 3 zero pixel by original pixel
2. Size is typically (2 x 2) but it can be any integer
3. Up-sampling:
  - Improves resolution for better features
  - Required in adder and concatenate blocks to add from previous outputs of higher resolutions
4. Bilinear up-sampling can be used for better quality at the cost of increase computations



# Fully connected neural network



# POPULAR FRAMEWORKS IN PYTHON FOR DEEP LEARNING/MACHINE LEARNING


---




**Caffe2**

PYTORCH

Caffe



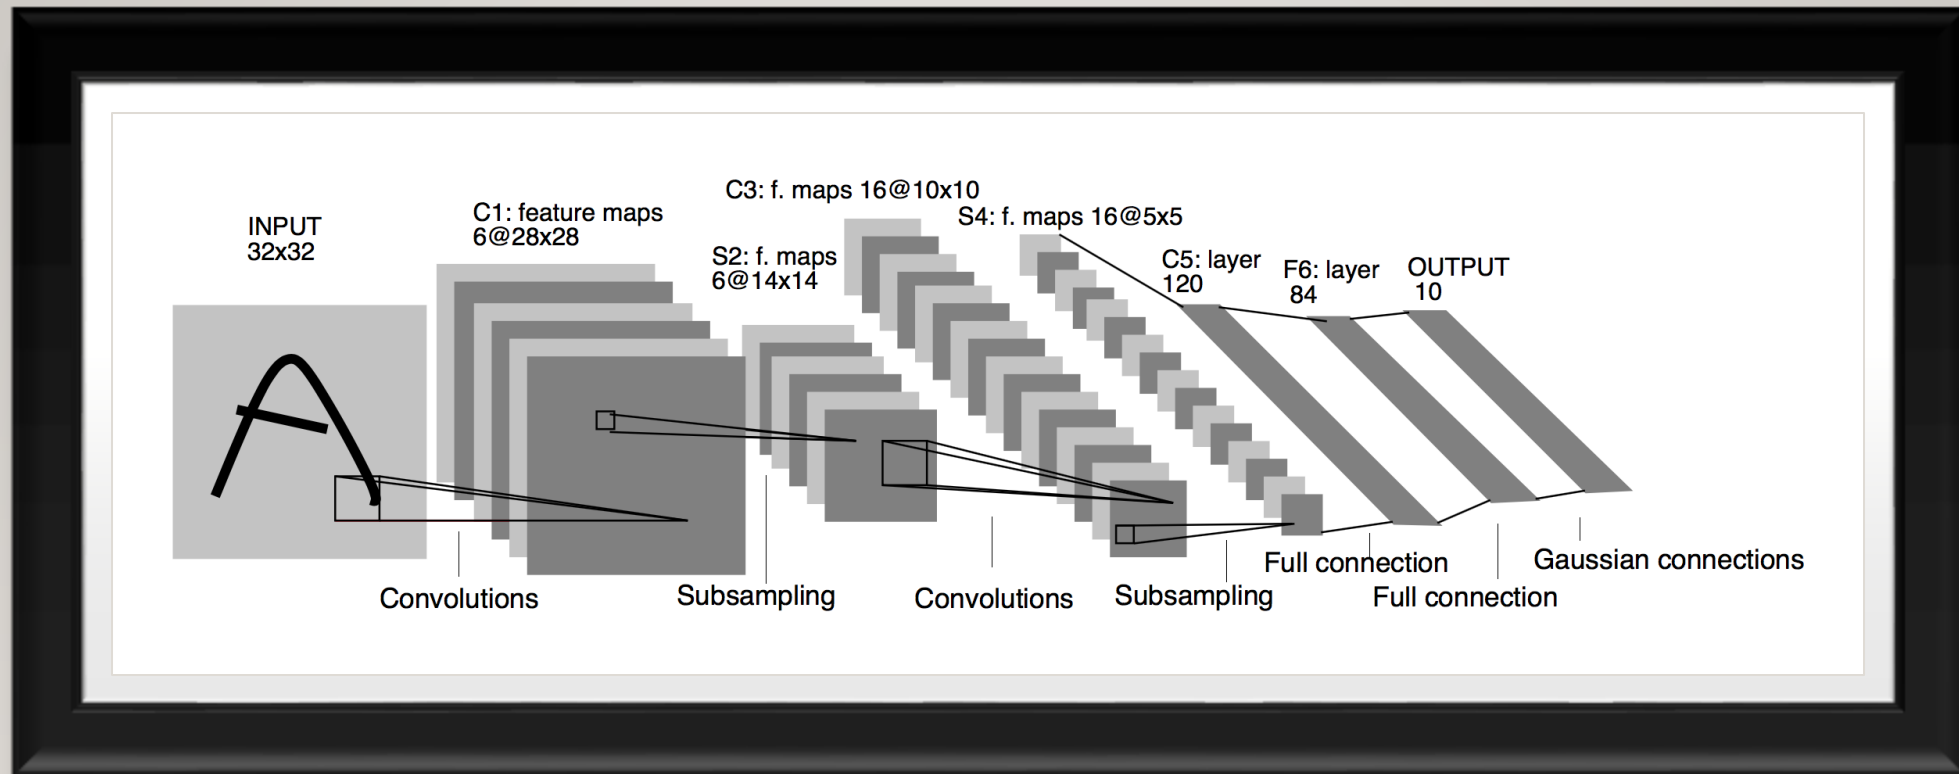
**Keras**



# IMAGENET COMPITITION

---



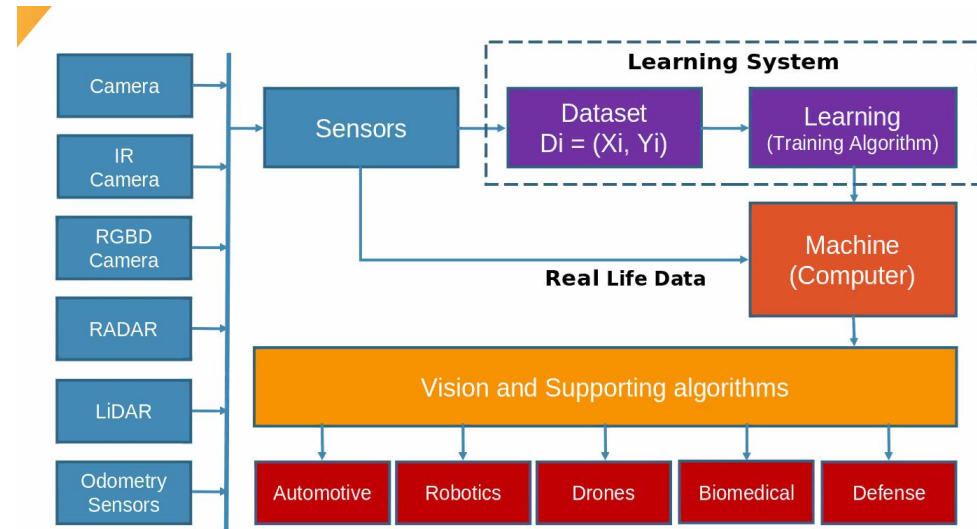


# LENET ARCHITECTURE FOR DIGIT RECOGNITION

## HOW TO APPROACH A IMAGE CLASSIFICATION PROBLEM?

---

- Train Dataset
- Training algorithm
- Loss function
- Optimizer
- Vision algorithm
- Test Dataset



# POPULAR DATASETS FOR COMPUTER VISION

---

- MNIST
- IMAGENET
- COCO
- CIFAR-10
- VISUALQA

# PREPARING DATASET USING KERAS

---

- We choose MNIST dataset for digit recognition.
- We divide the training dataset into two sets
  - Train Dataset
  - Test Dataset

```
def load_from_mnist(self):  
    mnist=keras.datasets.mnist  
    (x_train, x_train_label),(x_test, x_test_label) = mnist.load_data()  
    x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)  
    x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)  
    return x_test,x_test_label,x_train,x_train_label
```



# WRITING A SEQUENTIAL MODEL

---

- We create a sequential model and add layer by layer.

```
def gen_model(self):
    self.model.add(keras.layers.Conv2D(filters=6, kernel_size=5, strides=(1, 1), padding="same", input_shape=(28, 28, 1),
                                         data_format="channels_last", activation="relu"))
    self.model.add(keras.layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
    self.model.add(keras.layers.Conv2D(filters=16, kernel_size=5, strides=(1, 1), padding="valid", activation="relu"))
    self.model.add(keras.layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
    self.model.add(keras.layers.Flatten())
    self.model.add(keras.layers.Dense(120, input_shape=(400,), activation='relu'))
    self.model.add(keras.layers.Dense(84, input_shape=(120,), activation='relu'))
    self.model.add(keras.layers.Dense(units=10, activation='softmax'))
```

# TRAIN MODEL

---

- We need to select the optimizer and loss function based on our requirements.

```
def train_model(self):  
    self.model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])  
    # self.model.compile(loss='mean_squared_error', optimizer='sgd', metrics=['accuracy'])  
    self.model.fit(self.x_train, self.x_train_label, epochs=self.no_of_epochs, batch_size=self.batch_size)
```

```
def predict_output(self):  
    # loss_and_metrics = model.evaluate(x_test, x_test_label, batch_size=32)  
    classes = self.model.predict(self.x_test, batch_size=self.batch_size)  
    classes = classes.argmax(axis=1)  
    for row in range(0, 3):  
        plt.title("label=%s" % classes[row])  
        plt.imshow(np.reshape(self.x_test[row], (28, 28)), cmap='gray')  
        plt.show()
```

## EVALUATE/TEST MODEL

---

# ANY QUESTIONS?

---

Contact:

G Venkata Durga Rao

9000486749

venkatadj@gmail.com

